



# Serverless Data Processing with Dataflow

# Agenda

---

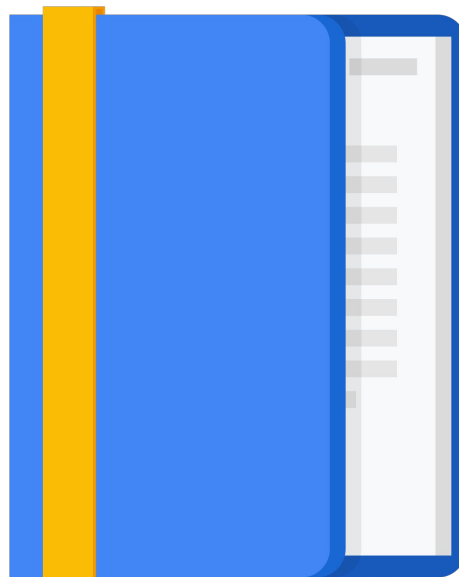
## Cloud Dataflow

Why customers value Dataflow

Dataflow Pipelines

Dataflow Templates

Dataflow SQL



# Google Cloud processing options (1)



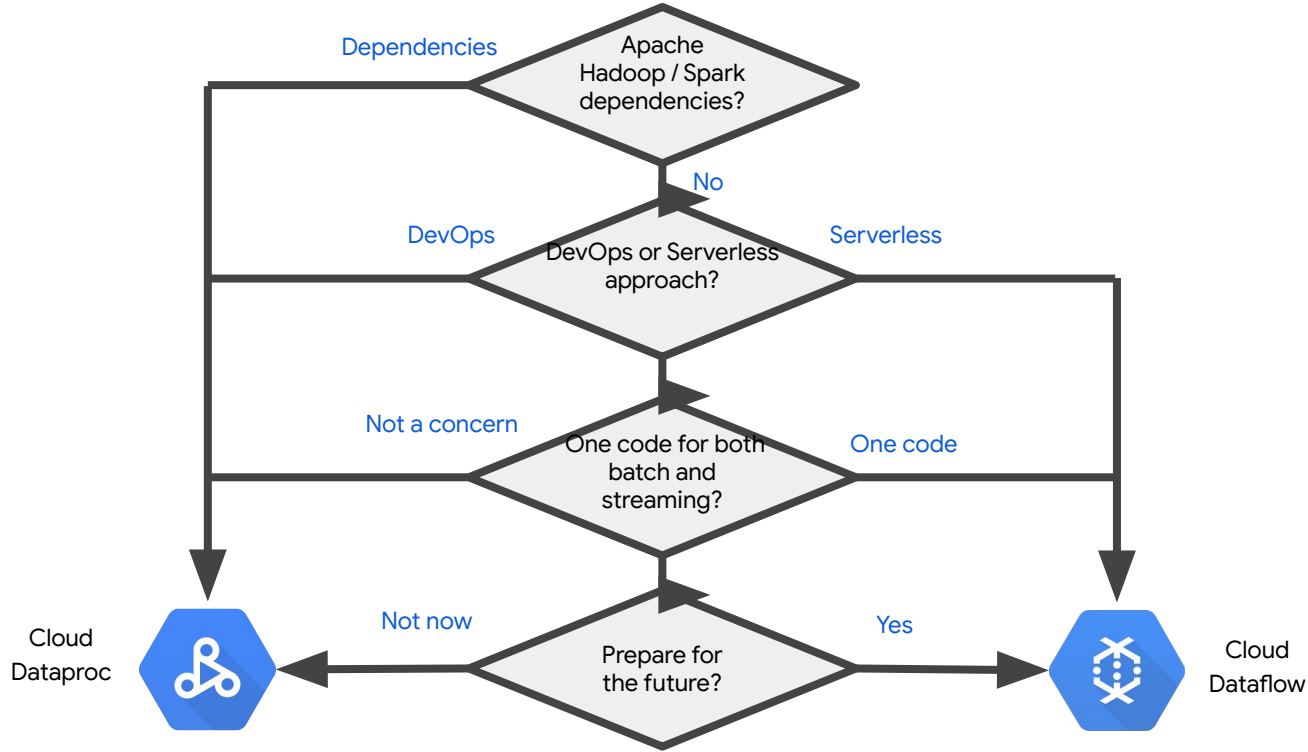
Cloud Dataflow



Cloud Dataproc

Recommended for:	New data processing pipelines, unified batch and streaming	Existing Hadoop/Spark applications, machine learning/data science ecosystem, large-batch jobs, preemptible VMs
Fully-managed:	Yes	No
Auto-scaling:	Yes, transform-by-transform (adaptive)	Yes, based on cluster utilization (reactive)
Expertise:	Apache Beam	Hadoop, Hive, Pig, Apache Big Data ecosystem, Spark, Flink, Presto, Druid

# Choosing between Cloud Dataflow and Cloud Dataproc



# Cloud Dataflow

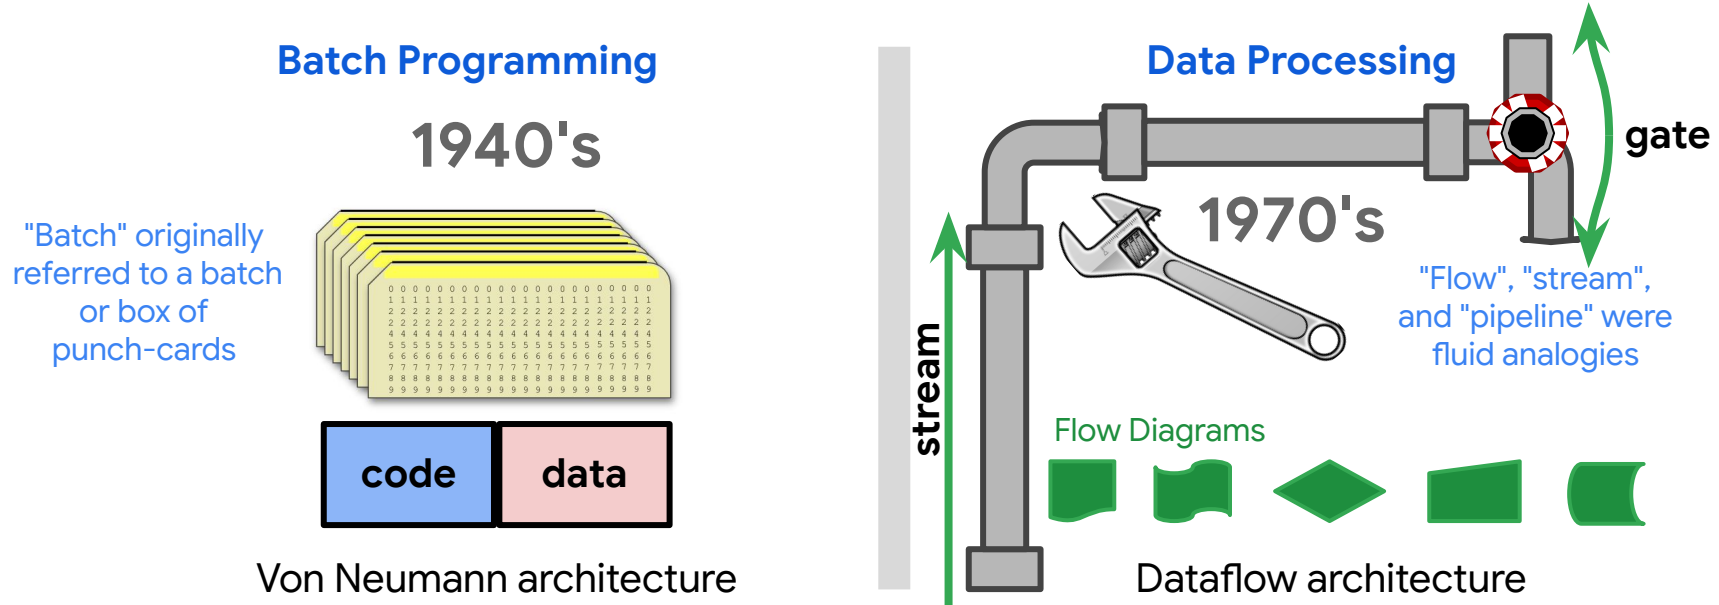


Cloud  
Dataflow

Qualities that Cloud Dataflow  
contributes to Data Engineering  
solutions:

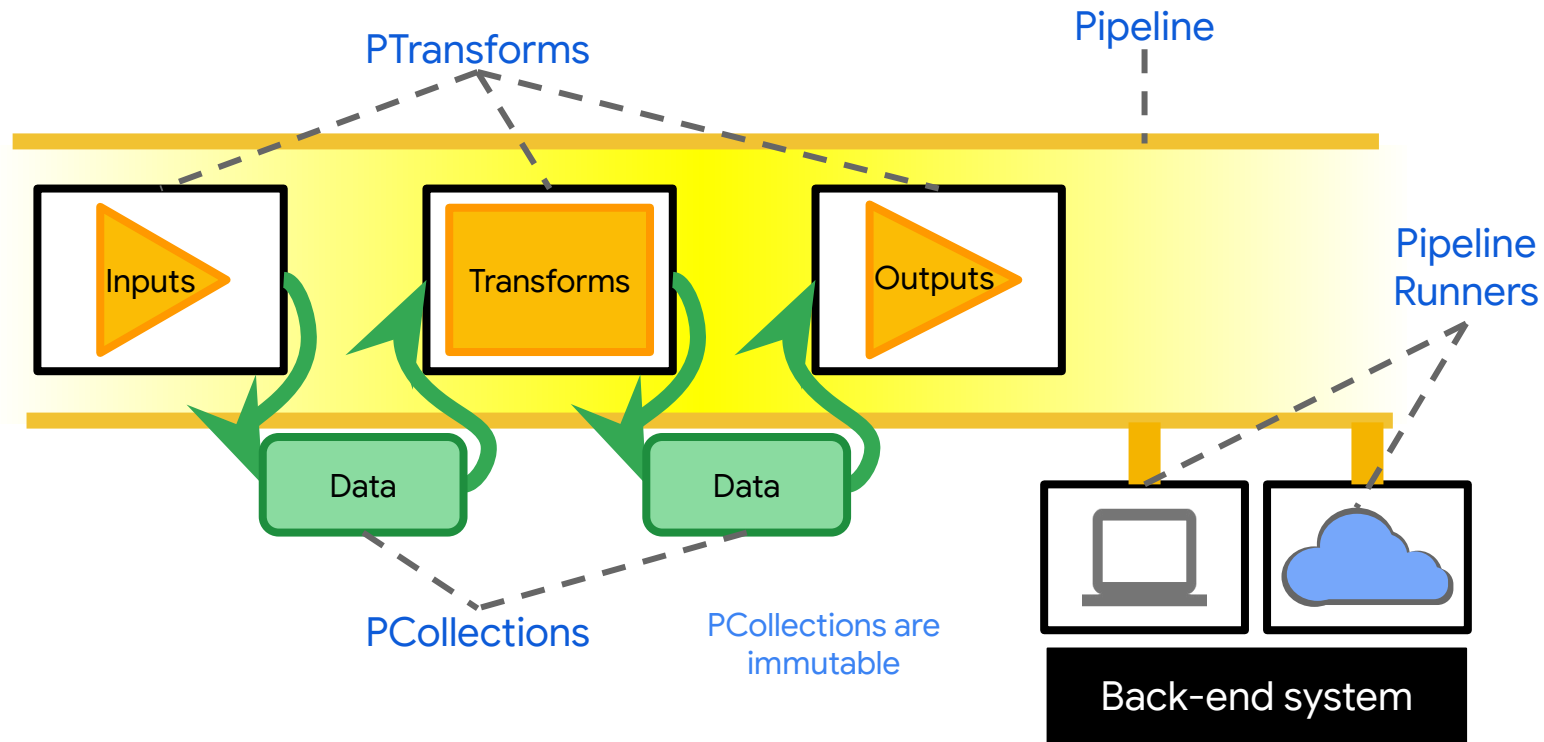
Scalability  
Low latency

# Batch programming and data processing used to be two very separate and different things

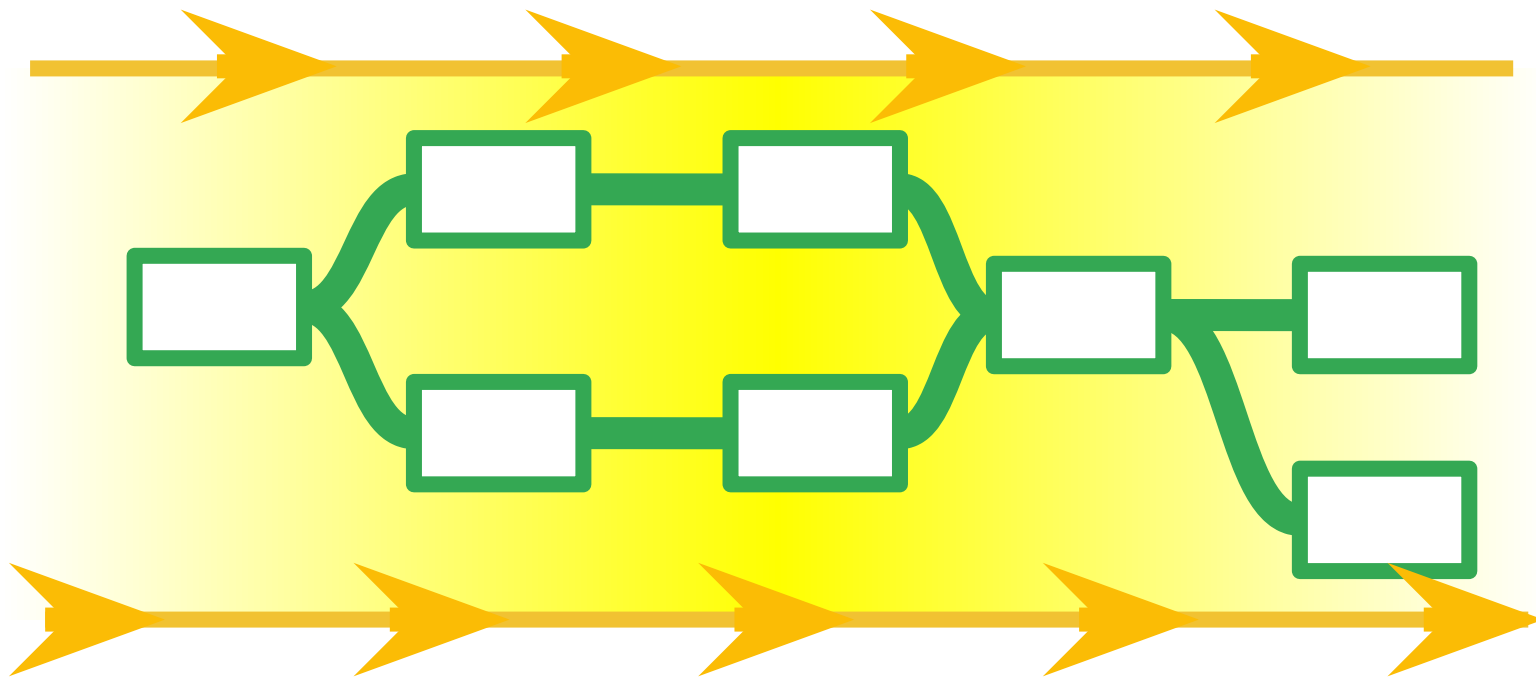


*Different tools, different platforms, different concepts, different methods.*

# Apache BEAM = Batch + strEAM

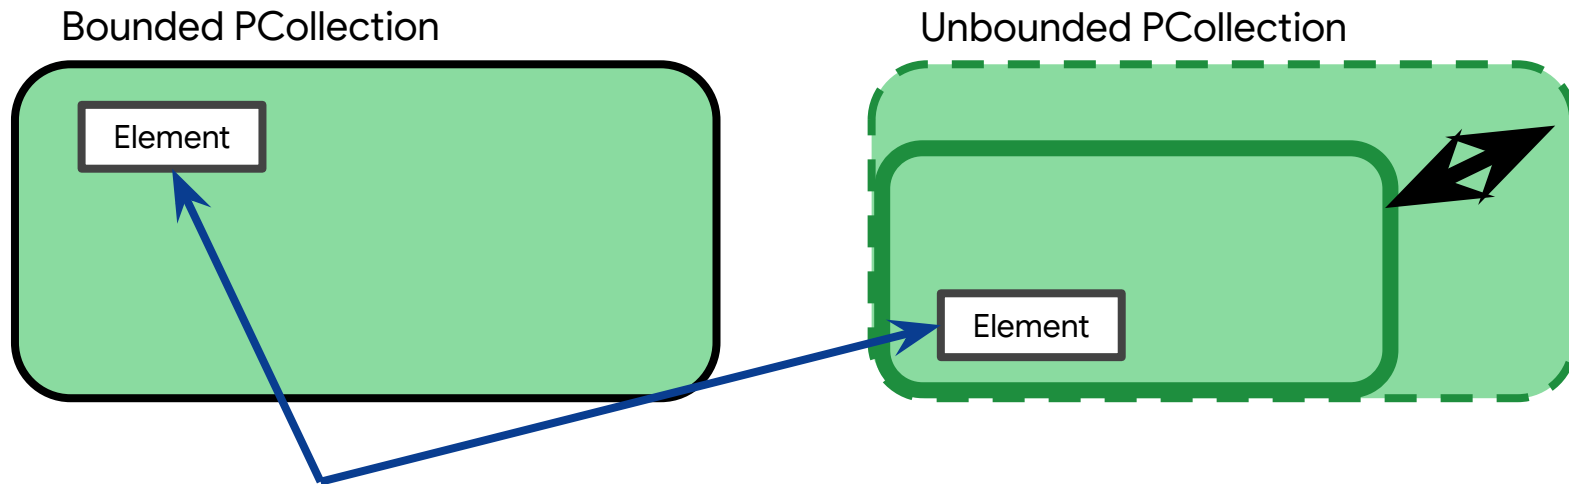


A Cloud Dataflow pipeline is a directed graph of steps





# A PCollection represents batch or stream data



All data types are stored  
as serialized byte strings

Note: Bounded means the data has a fixed size not that the PCollection size is limited. A PCollection can be any size and be distributed across many workers.

# Agenda

---

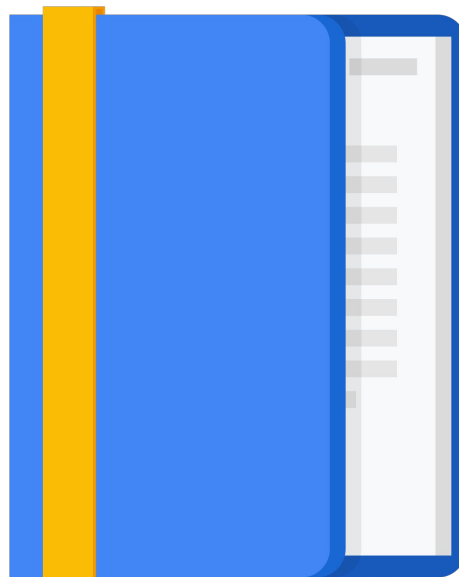
Cloud Dataflow

Why customers value Dataflow

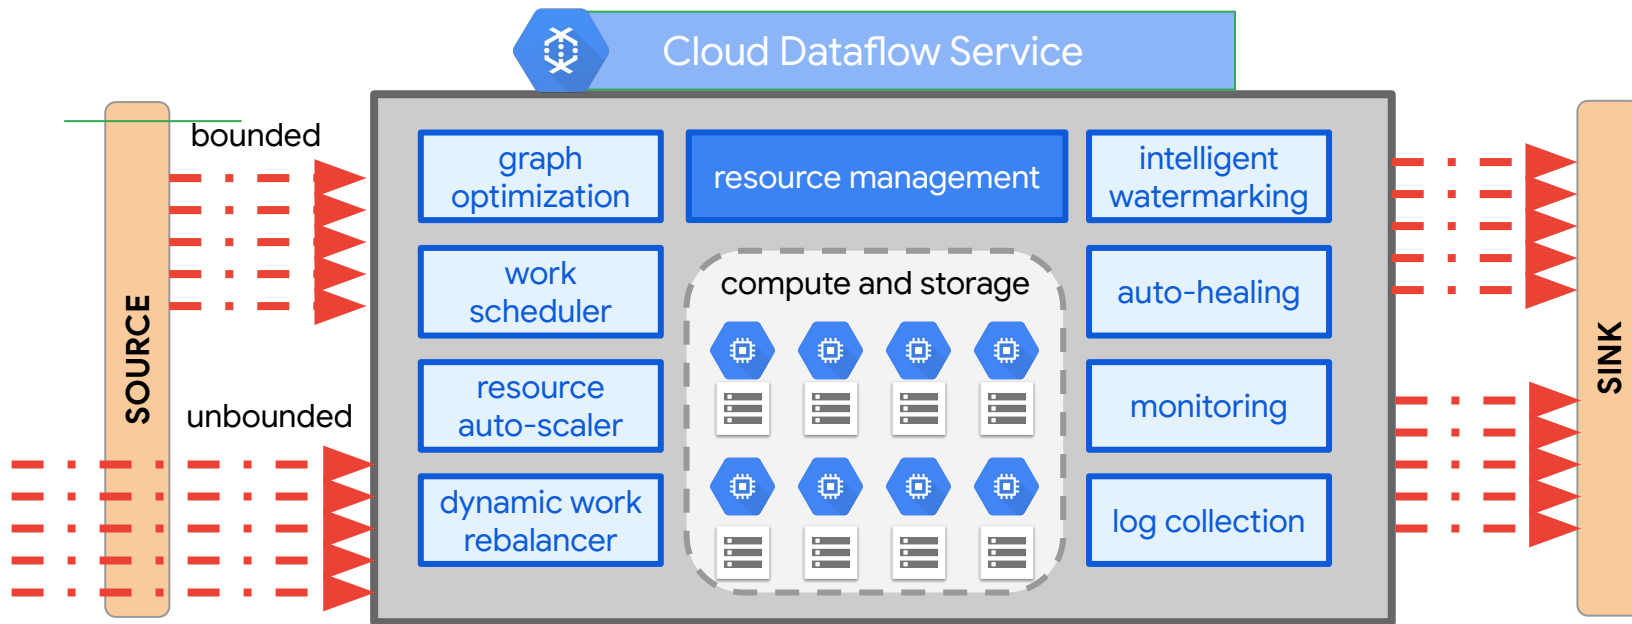
Dataflow Pipelines

Dataflow Templates

Dataflow SQL

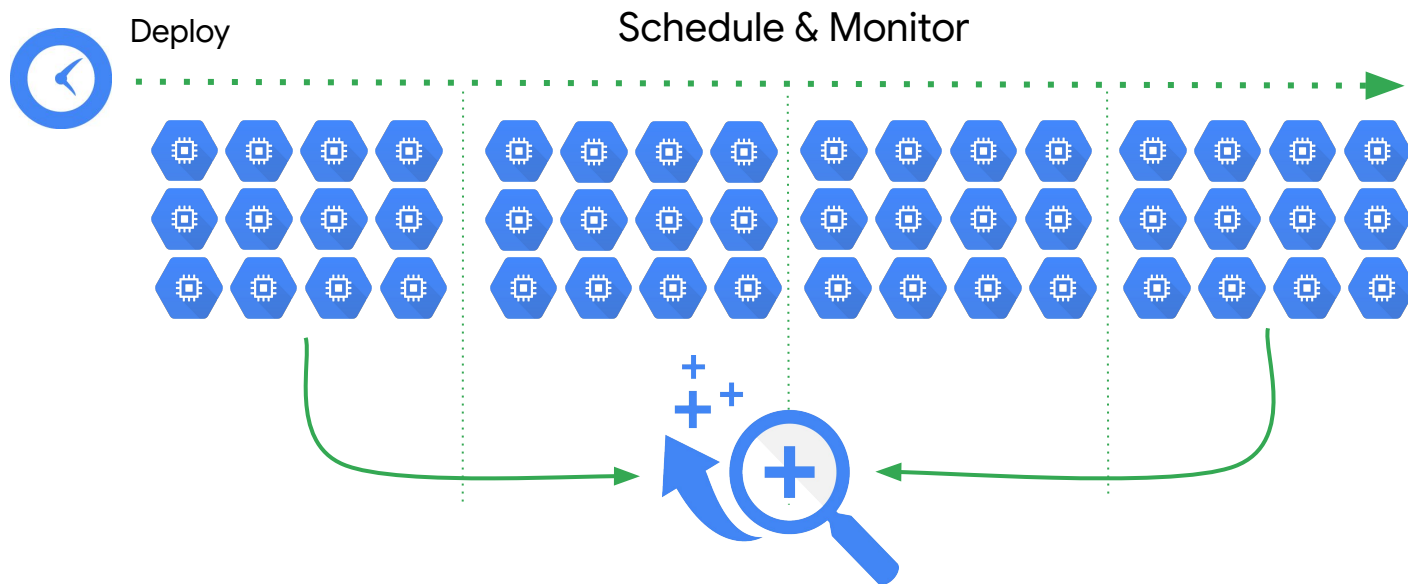


# How does Cloud Dataflow work?

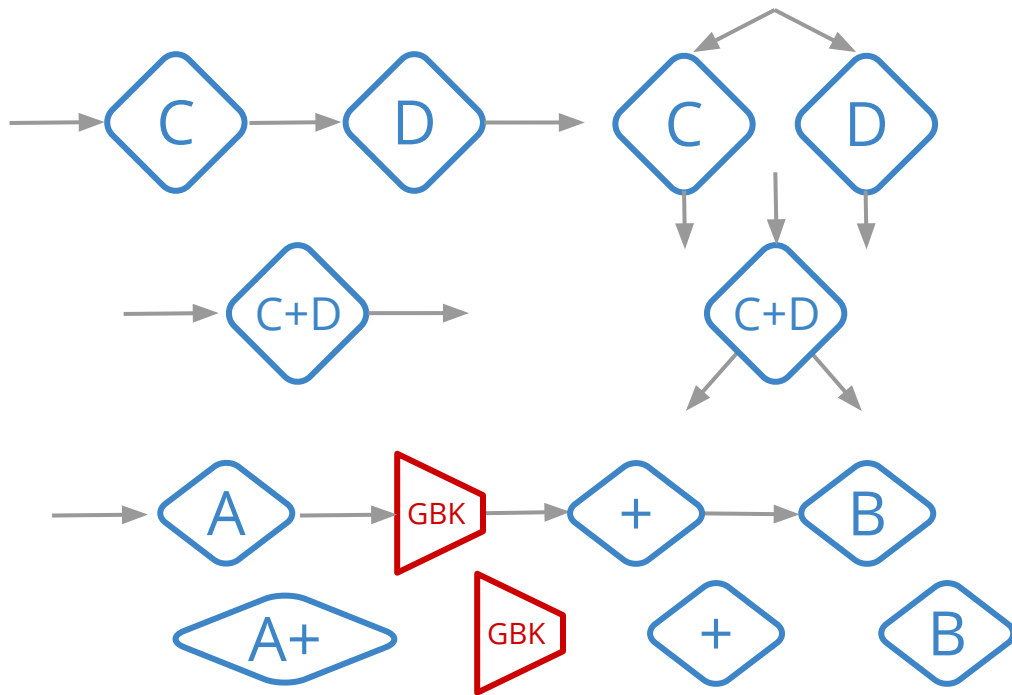


Cloud Dataflow constantly rebalances the work.

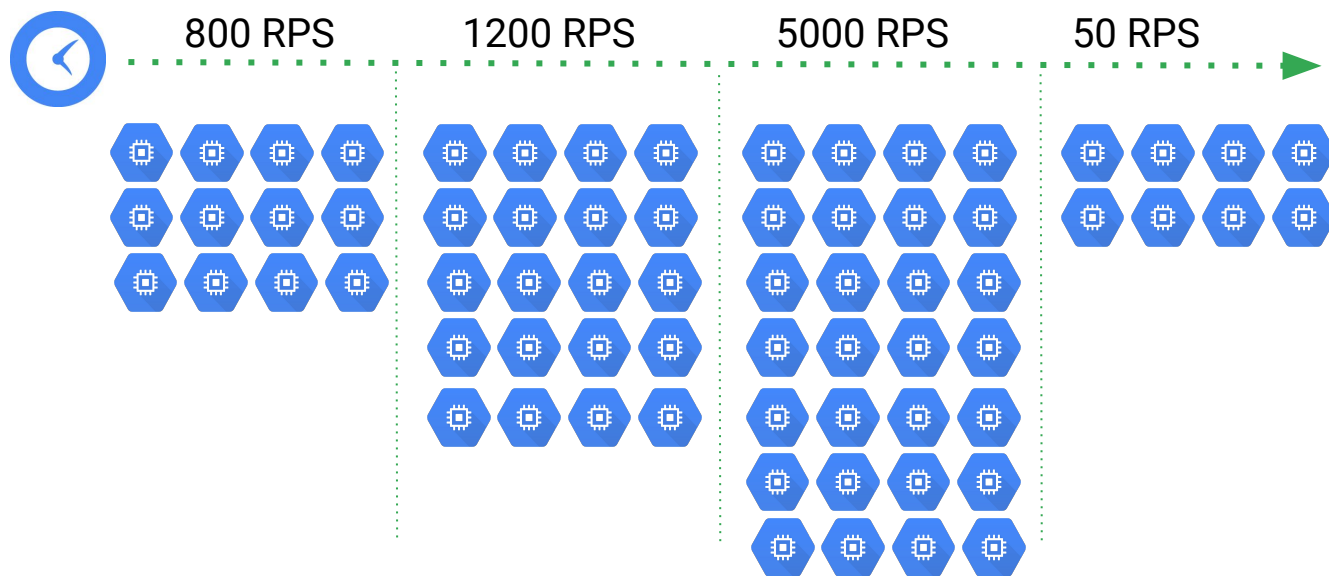
# Why customers value Cloud Dataflow: Fully-managed and auto-configured



# Why customers value Cloud Dataflow: Graph is optimized for best execution path



# Why customers value Cloud Dataflow: Autoscaling mid-job



# Why customers value Cloud Dataflow: Dynamic work rebalancing mid-job



# Why customers value Cloud Dataflow: Strong streaming semantics



Exactly once aggregations



Rich time tracking



Good integration with other GCP services



# Agenda

---

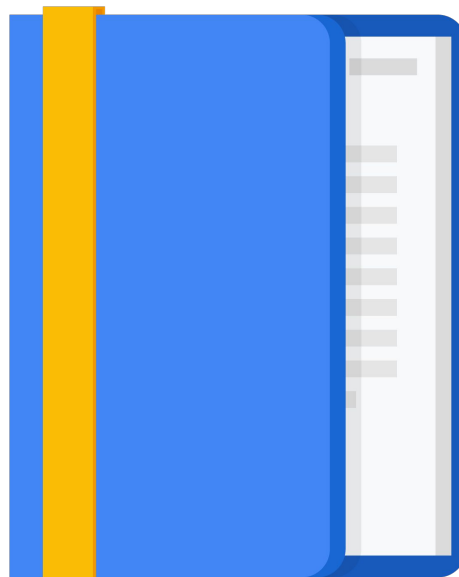
Cloud Dataflow

Why customers value Dataflow

Dataflow Pipelines

Dataflow Templates

Dataflow SQL



# How to construct a simple pipeline



```
PCollection_out = (PCollection_in | PTransform_1  
                  | PTransform_2  
                  | PTransform_3 )
```

Python

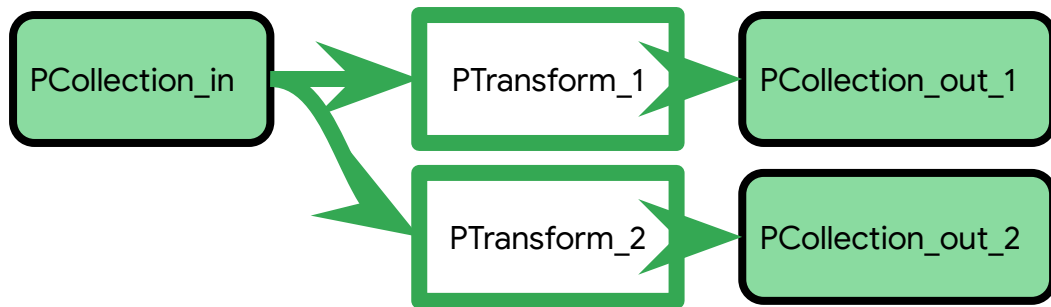
Python overloads  
the pipe operator

Java

Java uses the  
.apply method

```
PCollection_out = PCollection_in.apply(PTransform_1)  
                        .apply(PTransform_2)  
                        .apply(PTransform_3)
```

# How to construct a branching pipeline



```
PCollection_out_1 = PCollection_in | PTTransform_1  
PCollection_out_2 = PCollection_in | PTTransform_2
```

Python

Java

```
PCollection_out_1 = PCollection_in.apply(PTTransform_1)  
PCollection_out_2 = PCollection_in.apply(PTTransform_2)
```

# A Pipeline is a directed graph of steps

```
import apache_beam as beam

if __name__ == '__main__':
    with beam.Pipeline(argv=sys.argv) as p:

        (p
         | beam.io.ReadFromText('gs://...')
         | beam.FlatMap(lambda line:
count_words(line))
         | beam.io.WriteToText('gs://...')
        )

# end of with-clause: runs, stops the pipeline
```

Python

Create a pipeline  
parameterized by  
command line flags

Read input

Apply transform

Write output

# Run a pipeline on Cloud Dataflow

```
import apache_beam as beam
```

Python

```
options = {'project': <project>,  
           'runner': 'DataflowRunner',  
           'region': <region>,  
           'setup_file': <setup.py file>}  
pipeline_options =  
beam.pipeline.PipelineOptions(flags=[], **options)  
pipeline = beam.Pipeline(options = pipeline_options)
```

← — — — — — Where to run

← — — — — — This creates the pipeline

# Pipeline Execution using DataflowRunner

## Run local

```
python ./grep.py
```

## Run on cloud

```
python ./grep.py \  
  --project=$PROJECT \  
  --job_name=myjob \  
  --staging_location=gs://$BUCKET/staging/ \  
  --temp_location=gs://$BUCKET/tmp/ \  
  --runner=DataflowRunner
```

## Designing Pipelines

- **Input and Output**
- PTransforms

# Read data from local file system, Cloud Storage, Cloud Pub/Sub, BigQuery, ...

```
with beam.Pipeline(options=pipeline_options) as p:
```

## Read from Cloud Storage (returns a string)

```
lines = p | beam.io.ReadFromText("gs://.../input-*.csv.gz")
```

## Read from Cloud Pub/Sub (returns a string)

```
lines = p | beam.io.ReadStringsFromPubSub(topic=known_args.input_topic)
```

## Read from BigQuery (returns rows)

```
query = "SELECT x, y, z FROM [project:dataset.tablename]"  
BQ_source = beam.io.BigQuerySource(query = <query>, use_standard_sql=True)  
BQ_data = pipeline | beam.io.Read(BQ_source)
```

Setup

Read



# Write to a BigQuery table

## Establish reference to BigQuery table

```
from apache_beam.io.gcp.internal.clients import bigquery

table_spec = bigquery.TableReference(
    projectId='clouddataflow-readonly',
    datasetId='samples',
    tableId='weather_stations')
```

## Write to BigQuery table


```
p | beam.io.WriteToBigQuery(
    table_spec,
    schema=table_schema,
    write_disposition=beam.io.BigQueryDisposition.WRITE_TRUNCATE,
    create_disposition=beam.io.BigQueryDisposition.CREATE_IF_NEEDED)
```

# Create a PCollection from in-memory data

```
city_zip_list = [  
    ('Lexington', '40513'),  
    ('Nashville', '37027'),  
    ('Lexington', '40502'),  
    ('Seattle', '98125'),  
    ('Mountain View', '94041'),  
    ('Seattle', '98133'),  
    ('Lexington', '40591'),  
    ('Mountain View', '94085'),  
]  
citycodes = p | 'CreateCityCodes' >> beam.Create(city_zip_list)
```

Python

This is the display name  
of the pipeline step



PCollection

## Designing Pipelines

- Input and Output
- **PTransforms**

# Map and FlatMap


Use Map for 1:1 relationship between input and output

```
'WordLengths' >> beam.Map( lambda word: (word, len(word)) )
```

Map (fn) uses a callable fn to do a one-to-one transformation.


Use FlatMap for non 1:1 relationships, usually with a generator

```
def my_grep(line, term):  
    if term in line:  
        yield line  
  
'Grep' >> beam.FlatMap( lambda line: my_grep(line, searchTerm) )
```



FlatMap is similar to Map, but fn returns an iterable of zero or more elements.  
The iterables are flattened into one PCollection.

# ParDo implements parallel processing



ParDo acts on one item at a time in the PCollection  
Multiple instances of class on many machines  
Should not contain any state

## Uses:

- Filtering a data set, choosing which elements to output.
- Formatting or type-converting each element in a data set.
- Extracting parts of each element in a data set.
- Performing computations on each element in a data set.

# ParDo requires code passed as a DoFn object

```
words = ...
```

```
class ComputeWordLengthFn(beam.DoFn):  
    def process(self, element):  
        return [len(element)]
```

```
word_lengths = words | beam.ParDo(ComputeWordLengthFn())
```

Python

The input is a PCollection of strings.

The DoFn to perform on each element in the input PCollection.

The output is a PCollection of integers.

Apply a ParDo to the PCollection "words" to compute lengths for each word.

## ParDo method can emit multiple variables

```
results = (words | beam.ParDo(ProcessWords(), cutoff_length=2, marker='x')
    .with_outputs('above_cutoff_lengths', 'marked strings',
main='below_cutoff_strings'))

below  = results.below_cutoff_strings
above  = results.above_cutoff_lengths
marked = results['marked strings']
```



---

## A Simple Dataflow Pipeline (Python/Java)


### Objectives

- Open Dataflow project
- Pipeline filtering
- Execute the pipeline locally and on the cloud



# GroupByKey explicitly shuffles key-values pairs


```
cityAndZipcodes = p | beam.Map(lambda fields : (fields[0], fields[1]))  
grouped = cityAndZipCodes | beam.GroupByKey()
```



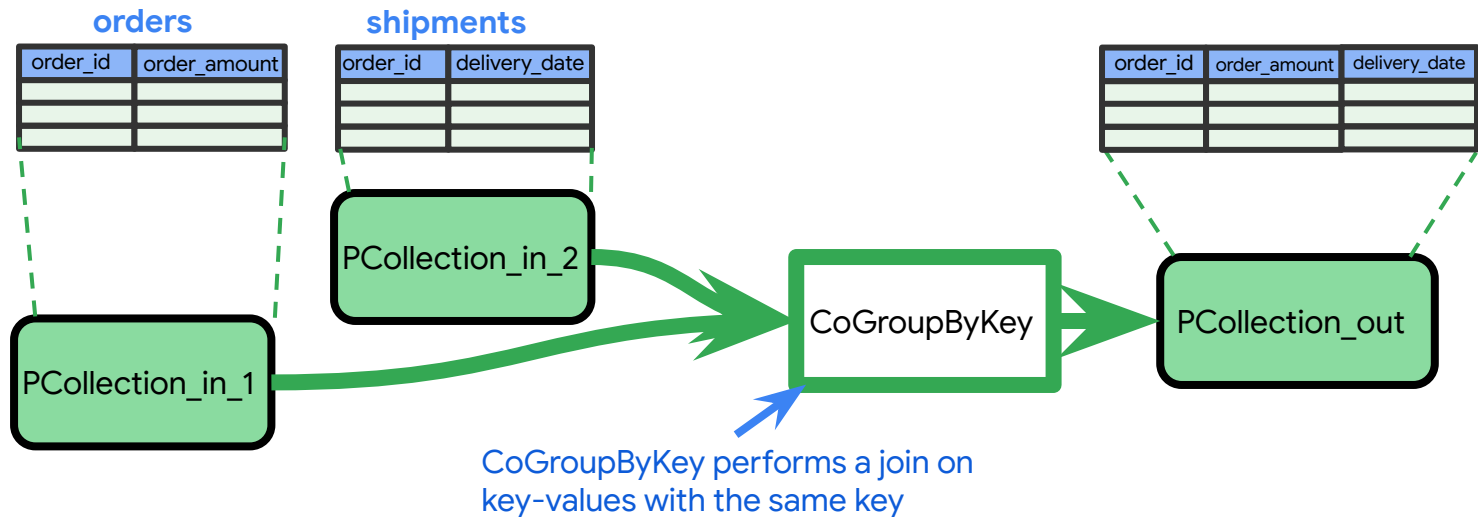
```
Lexington, 40513  
Nashville, 37027  
Lexington, 40502  
Seattle, 98125  
Mountain View, 94041  
Seattle, 98133  
Lexington, 40591  
Mountain View, 94085
```

```
Lexington, [40513, 40502, 40592]  
Nashville, [37027]  
Seattle, [98125, 98133]  
Mountain View, [94041, 94085]
```

# Data skew makes grouping less efficient at scale



# CoGroupByKey joins two or more key-value pairs



```
results = ({'orders': orders, 'shipments': shipments}
           | beam.CoGroupByKey())
```


# Combine (reduce) a PCollection

## Applied to a PCollection of values


```
totalAmount = salesAmounts | CombineGlobally(sum)
```

## Applied to a grouped Key-Value pair

```
totalSalesPerPerson = salesRecords | CombinePerKey(sum)
```



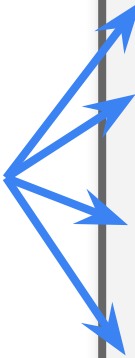
Each element of salesRecords  
is a tuple:  
(salesPerson, salesAmount)



Pre-built combine functions  
for many common numeric  
combination operations such  
as sum, mean, min, and max

# CombineFn works by overriding existing operations

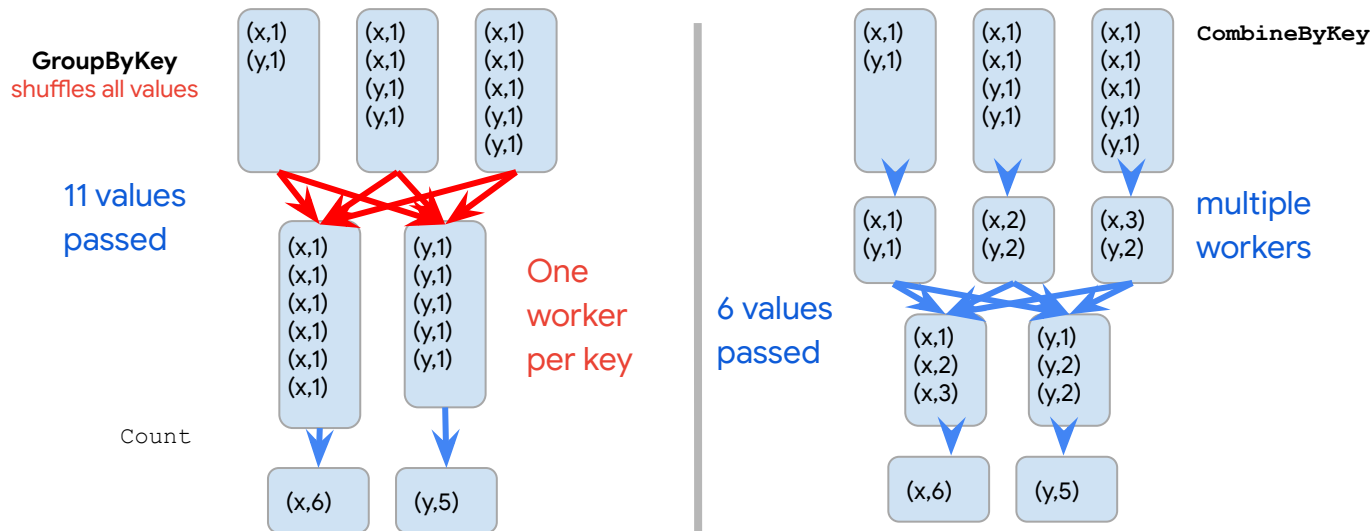
You must provide four operations by overriding the corresponding methods



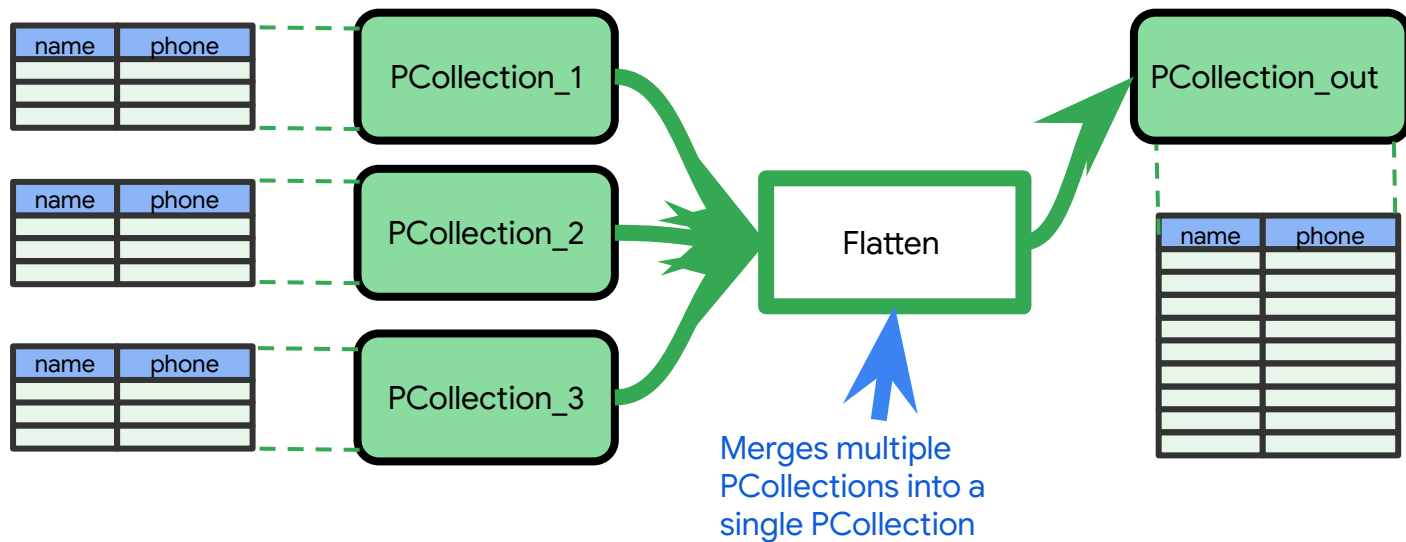
```
class AverageFn(beam.CombineFn):  
  
    def create_accumulator(self):  
        return (0.0, 0)  
  
    def add_input(self, sum_count, input):  
        (sum, count) = sum_count  
        return sum + input, count + 1  
  
    def merge_accumulators(self, accumulators):  
        sums, counts = zip(*accumulators)  
        return sum(sums), sum(counts)  
  
    def extract_output(self, sum_count):  
        (sum, count) = sum_count  
        return sum / count if count else float('NaN')
```

```
pc = ...  
average = pc | beam.CombineGlobally(AverageFn())
```

# Combine is more efficient than GroupByKey

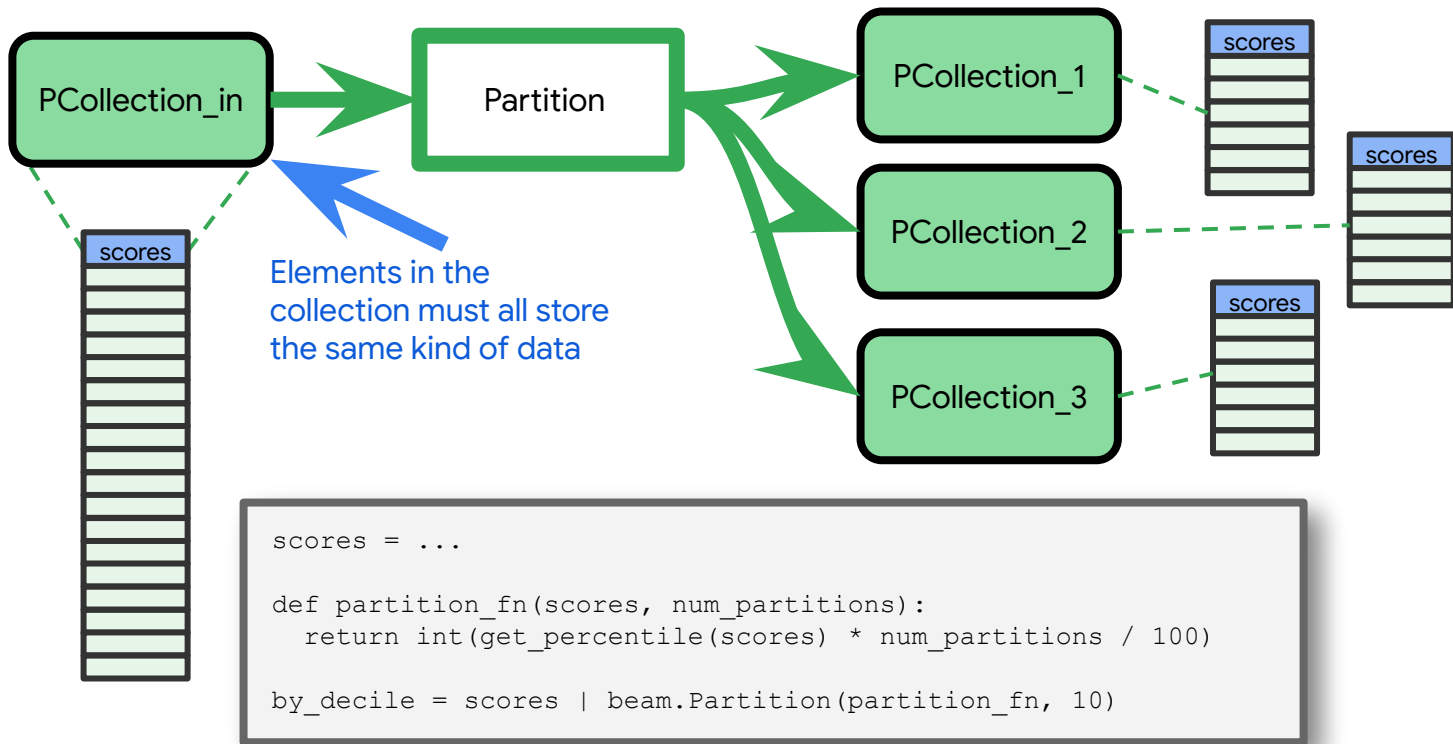


# Flatten merges identical PCollections



```
merged = ((pcoll1, pcoll2, pcoll3) | beam.Flatten())
```

# Partition splits PCollections into smaller PCollections







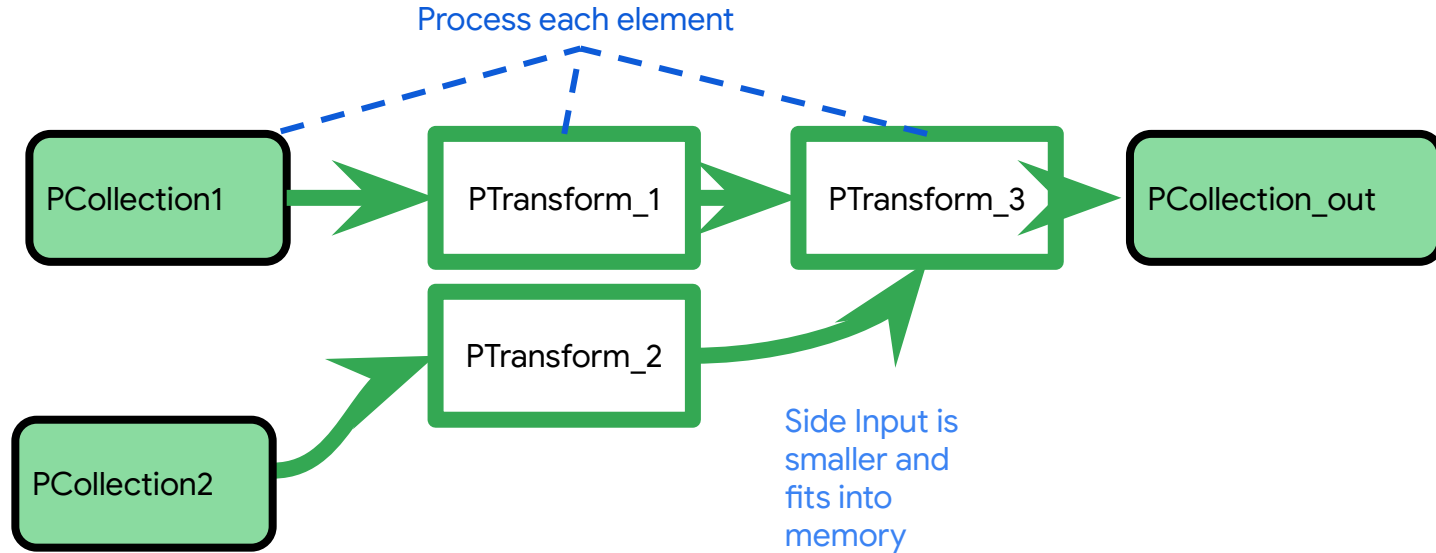
---

## MapReduce in Dataflow (Python/Java)

### Objectives

- Identify Map and Reduce operations
- Execute the pipeline
- Use command line parameters

# Use side inputs to inject additional runtime data



# How side inputs work

```
words = ...

def filter_using_length(word, lower_bound, upper_bound=float('inf')):
    if lower_bound <= len(word) <= upper_bound:
        yield word

small_words = words | 'small' >> beam.FlatMap(filter_using_length, 0, 3)

avg_word_len = (words
                | beam.Map(len)
                | beam.CombineGlobally(beam.combiners.MeanCombineFn()))

larger_than_average = (words | 'large' >> beam.FlatMap(
    filter_using_length,
    lower_bound=pvalue.AsSingleton(avg_word_len)))
```

Side input



---

## Side Inputs (Python/Java)


### Objectives

- Try out a BigQuery query
- Explore the pipeline code
- Execute the pipeline

# Processing Time-series data using Windowing

# Every PCollection is processed within a Window


## Bounded PCollection



- 2 In Bounded PCollections, commonly the Elements are all marked as occurring at the same time. (Example: TextIO does this.) So the global window basically ignores the timing information.

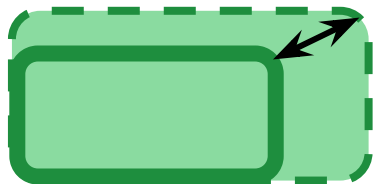


- 1 The default window is called the global window, it starts when the data is input and ends when the last element in the collection is processed.



# The global window is not very useful for an unbounded PCollection

## Unbounded PCollection




1

The timing associated with the elements in an Unbounded PCollection is usually important to processing the data.

3

The discussion about Unbounded PCollections and Windows will be continued in the course on Processing Streaming Data.



2

An Unbounded PCollection has no defined end or last element. So it can never perform the completion step.

This is particularly important for **GroupByKey** and **Combine**, which perform the shuffle after 'end'.

# Setting a single global window for a PCollection.

## Single global window

```
from apache_beam import window
session_windowed_items = (
    items | 'window' >> beam.WindowInto(window.GlobalWindows()))
```

Python

This is the default.

This code illustrates how you could explicitly set it.

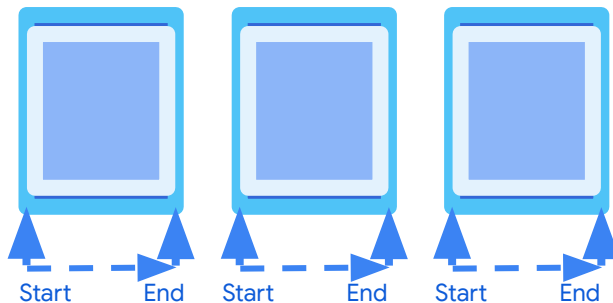


# Time-based Windows can be useful for processing time-series data



1

You may have to prepare the date-timestamp. In this example, the dts of the data (log writing time) becomes the element time. Now the elements have different times from one another.



2

Using time based windowing the data is processed in groups.

In the example, each group gets its own average.

3

There are different kinds of windowing.

Shown is "Fixed" There is also "Sliding" and "Session".

# Using Windowing with Batch (group by time)

```
lines = p | 'Create' >> beam.io.ReadFromText('access.log')
windowed_counts = (
    lines
    | 'Timestamp' >> beam.Map(lambda x: beam.window.TimestampedValue(x, extract_timestamp(x)))
    | 'Window' >> beam.WindowInto(beam.window.SlidingWindows(60, 30))
    | 'Count' >> (beam.CombineGlobally(beam.combiners.CountCombineFn()).without_defaults())
)
windowed_counts = windowed_counts | beam.ParDo(PrintWindowFn())
```

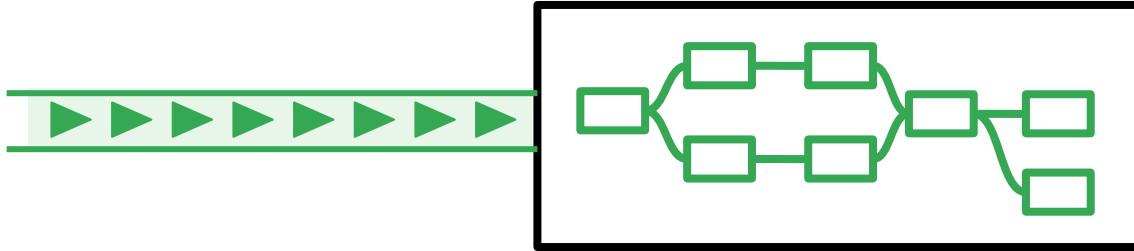
Python

## access.log (example)

```
131.108.5.17 - - [29/Apr/2019:04:53:15 -0800] "GET /view HTTP/1.1" 200 7352
131.108.5.17 - - [29/Apr/2019:05:21:35 -0800] "GET /view HTTP/1.1" 200 5253
```

Date Time Stamp

# Streaming data processing with Cloud Dataflow



Discussion of streaming continues in the  
Streaming Data Processing course.

# Agenda

---

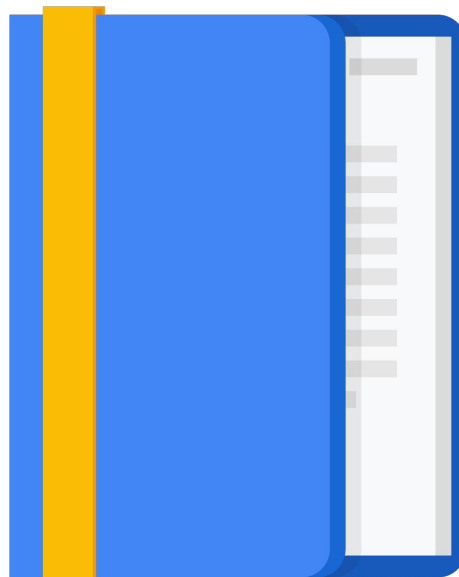
Cloud Dataflow

Why customers value Dataflow

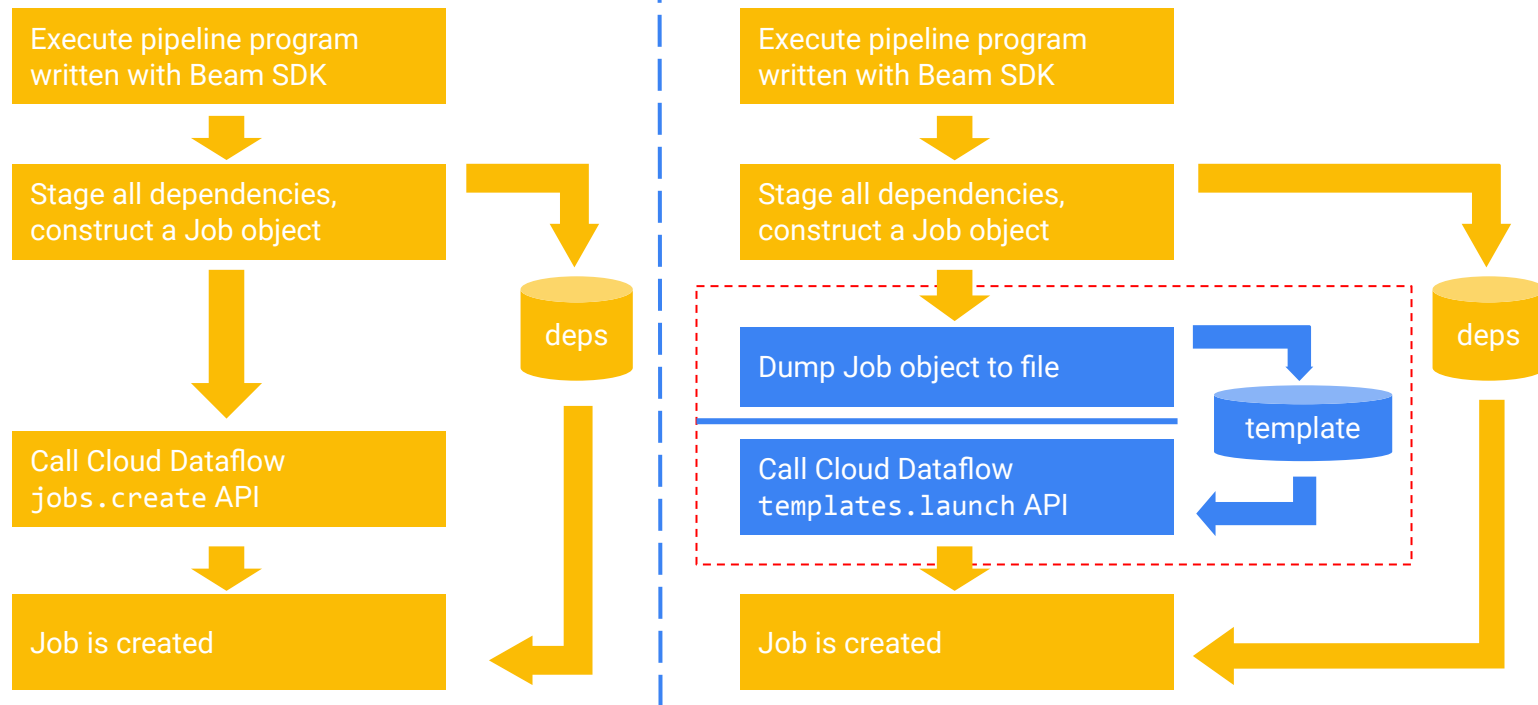
Dataflow Pipelines

Dataflow Templates

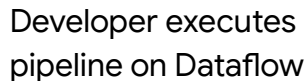
Dataflow SQL



# Cloud Dataflow templates enable the rapid deployment of standard job types



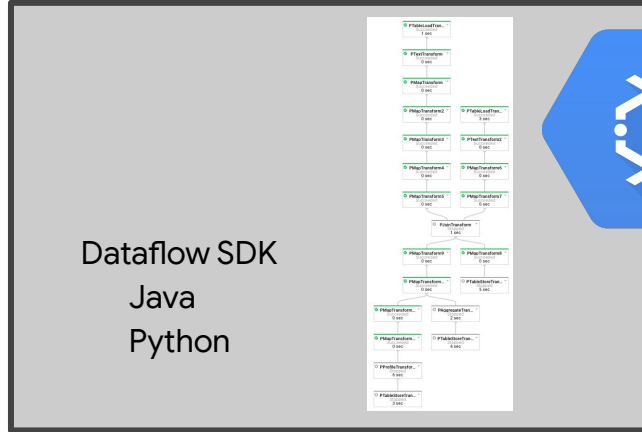
## Development environment



Developer or User  
submits source code to  
run Dataflow jobs

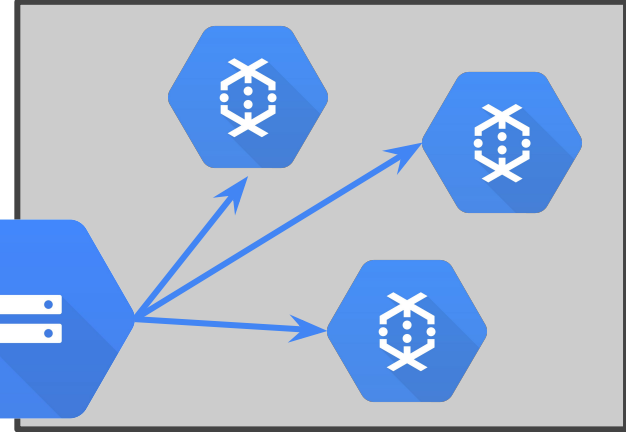
# Template workflow supports non-developer users

## Development environment



Developer creates pipeline in the development environment

## Production environment




Dataflow stores template in cloud storage

Users submit templates to run jobs

# Get started with Google-provided templates

Pre-written Cloud Dataflow pipelines for common data tasks that can be triggered with a single command or UI form.



## Target users

- App developers
- DB admins
- Analysts
- Data scientists
- Data engineers



## Exposure

- Through Google-provided Cloud Dataflow templates
- Embedded in other GCP products calling templates API

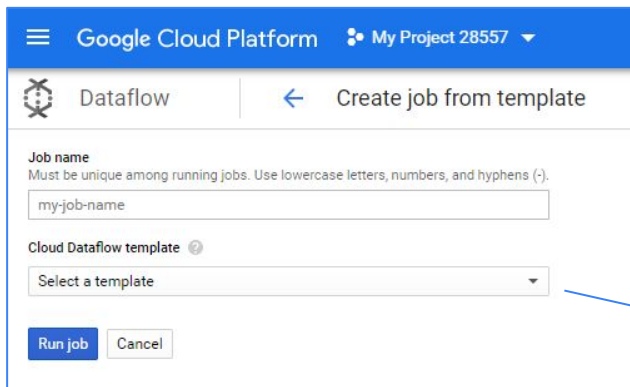


## Data Fusion

- Branded Google product
- UI pipeline builder
- Scheduler/orchestrator



# Execute templates with the GCP Console, gcloud command-line tool, or the REST API



The screenshot shows the Google Cloud Platform console interface for creating a Dataflow job. The top navigation bar includes the Google Cloud logo, 'Google Cloud Platform', and 'My Project 28557'. The main header has a 'Dataflow' tab and a 'Create job from template' button. Below this, the 'Job name' field is set to 'my-job-name' with a note: 'Must be unique among running jobs. Use lowercase letters, numbers, and hyphens (-)'. The 'Cloud Dataflow template' dropdown menu is open, showing 'Select a template'. At the bottom are 'Run job' and 'Cancel' buttons. A blue arrow points from the dropdown menu to a list of templates on the right.

```
gcloud dataflow jobs run \  
--gcs-location=gs://df-ts/latest/PubsubToBigQuery \  
--parameters inputTopic=X outputTable=Y
```

- 
- The list of templates is organized into two sections: 'Get Started' and 'Process Data in Bulk (batch)'. The 'Get Started' section includes 'Word Count' and a category 'Process Data Continuously (stream)' with templates like 'Cloud Pub/Sub Subscription to BigQuery' and 'Cloud Pub/Sub Topic to BigQuery'. The 'Process Data in Bulk (batch)' section includes templates like 'Text Files on Cloud Storage to Cloud Pub/Sub' and 'Text Files on Cloud Storage to BigQuery'.
- Get Started
    - Word Count
    - Process Data Continuously (stream)
      - Cloud Pub/Sub Subscription to BigQuery
      - Cloud Pub/Sub Topic to BigQuery
      - Cloud Pub/Sub to Text Files on Cloud Storage
      - Cloud Pub/Sub to Avro Files on Cloud Storage
      - Cloud Pub/Sub to Cloud Pub/Sub
      - Text Files on Cloud Storage to Cloud Pub/Sub
      - Text Files on Cloud Storage to BigQuery
      - Data Masking/Tokenization using Cloud DLP from GCS to BigQuery
  - Process Data in Bulk (batch)
    - Text Files on Cloud Storage to Cloud Pub/Sub
    - Text Files on Cloud Storage to BigQuery
    - Cloud Datastore to Text Files on Cloud Storage
    - Text Files on Cloud Storage to Cloud Datastore
    - Cloud Spanner to Text Files on Cloud Storage
    - Cloud Spanner to Avro Files on Cloud Storage
    - Avro Files on Cloud Storage to Cloud Spanner
    - Cloud BigTable to SequenceFile Files on Cloud Storage
    - SequenceFile Files on Cloud Storage to Cloud BigTable
    - Cloud Bigtable to Avro Files on Cloud Storage
    - Avro Files on Cloud Storage to Cloud Bigtable
    - Jdbc to BigQuery

# Google-provided templates documentation

How-to guides

All how-to guides

Installing the SDK

▶ Creating a pipeline

Specifying execution parameters

Deploying a pipeline

Using the monitoring UI

Using the command-line interface

Using Stackdriver Monitoring 📈

Logging pipeline messages

▶ Troubleshooting your pipeline

Updating an existing pipeline

Stopping a running pipeline

▼ Creating and executing templates

Overview

▼ Google-provided templates

Get started

Streaming templates

Batch templates

Utility templates

Creating templates

Executing templates

Migrating from MapReduce

Migrating from SDK 1.x for Java

▶ Configuring networking

Using Cloud Pub/Sub Seek

Using Flexible Resource Scheduling 📈

▶ Creating Cloud Dataflow SQL jobs 📈

Cloud Dataflow > Documentation

## Get started with Google-provided templates

Google provides a set of [open-source](#) Cloud Dataflow templates. For general information about templates, see the [Overview](#) page. To get started, use the [WordCount](#) template documented in the section below. See other Google-provided templates:

**Streaming templates** - Templates for processing data continuously:

- [Cloud Pub/Sub Subscription to BigQuery](#)
- [Cloud Pub/Sub Topic to BigQuery](#)
- [Cloud Pub/Sub to Cloud Pub/Sub](#)
- [Cloud Pub/Sub to Cloud Storage Avro](#)
- [Cloud Pub/Sub to Cloud Storage Text](#)
- [Cloud Storage Text to BigQuery \(Stream\)](#)
- [Cloud Storage Text to Cloud Pub/Sub \(Stream\)](#)
- [Data Masking/Tokenization using Cloud DLP from Cloud Storage to BigQuery \(Stream\)](#)

**Batch templates** - Templates for processing data in bulk:

- [Cloud Bigtable to Cloud Storage Avro](#)
- [Cloud Bigtable to Cloud Storage SequenceFiles](#)
- [Cloud Datastore to Cloud Storage Text](#)
- [Cloud Spanner to Cloud Storage Avro](#)
- [Cloud Spanner to Cloud Storage Text](#)
- [Cloud Storage Avro to Cloud Bigtable](#)

☆☆☆☆

[SEND FEEDBACK](#)

Contents

WordCount

# Use cases of Google-provided templates

- Code-free routine job launcher for data engineers
- Building block for import/export feature of other services on GCP
- OSS code base works as good knowledge base



Cloud Pub/Sub



Cloud Spanner



Cloud BigTable

## Which means now you can...

- Launch Dataflow jobs programmatically (via API).
- Launch Dataflow jobs instantaneously.
- Re-use Dataflow jobs
- Letting you customize the execution of your pipeline

# What if you want to create your own template?

- Doc: <https://cloud.google.com/dataflow/docs/templates/overview>
- Steps
  1. Modify pipeline options with ValueProviders.
  2. Generate template file.

```
mvn compile exec:java \  
-Dexec.mainClass=com.example.myclass \  
-Dexec.args="--runner=DataflowRunner \  
--project=[YOUR_PROJECT_ID] \  
--stagingLocation=gs://[YOUR_BUCKET_NAME]/staging \  
--output=gs://[YOUR_BUCKET_NAME]/output \  
--templateLocation=gs://[YOUR_BUCKET_NAME]/templates/MyTemplate"
```

3. Call it from API.

```
POST https://dataflow.googleapis.com/v1b3/projects/[YOUR_PROJECT_ID]/templates:launch?gcsPath=gs://[  
{  
  "jobName": "[JOB_NAME]",  
  "parameters": {  
    "inputFile": "gs://[YOUR_BUCKET_NAME]/input/my_input.txt",  
    "outputFile": "gs://[YOUR_BUCKET_NAME]/output/my_output"  
  },  
  "environment": {  
    "tempLocation": "gs://[YOUR_BUCKET_NAME]/temp",  
    "zone": "us-central1-f"  
  }  
}
```

# Templates require modifying parameters for runtime

```
class WordcountOptions(PipelineOptions):  
    @classmethod  
    def _add_argparse_args(cls, parser):  
        parser.add_value_provider_argument(  
            '--input',  
            default='gs://dataflow-samples/shakespeare/kinglear.txt',  
            help='Path of the file to read from')  
        parser.add_argument(  
            '--output',  
            required=True,  
            help='Output file to write results to.')  
    pipeline_options = PipelineOptions(['--output', 'some/output_path'])  
    p = beam.Pipeline(options=pipeline_options)  
  
    wordcount_options = pipeline_options.view_as(WordcountOptions)  
    lines = p | 'read' >> ReadFromText(wordcount_options.input)
```

Python

Run-time  
parameters

Non run-time  
parameters can stay

Runtime parameters  
must be modified

# Creating a template

- ValueProviders are passed down throughout the whole pipeline construction phase
- ValueProvider.get() only available in processElement()
  - Because it is fulfilled via API call

```
public interface SumIntOptions extends PipelineOptions {
    // New runtime parameter, specified by the --int
    // option at runtime.
    ValueProvider<Integer> getInt();
    void setInt(ValueProvider<Integer> value);
}

class MySumFn extends DoFn<Integer, Integer> {
    ValueProvider<Integer> mySumInteger;

    MySumFn(ValueProvider<Integer> sumInt) {
        // Store the value provider
        this.mySumInteger = sumInt;
    }

    @ProcessElement
    public void processElement(ProcessContext c) {
        // Get the value of the value provider and add it to
        // the element's value.
        c.output(c.element() + mySumInteger.get());
    }
}

public static void main(String[] args) {
    SumIntOptions options =
        PipelineOptionsFactory.fromArgs(args).withValidation()
            .as(SumIntOptions.class);
}
```

# Nested Value Providers

Sometimes we need to transform a value from what the user passes at Runtime to what a Source/Sink expects to consume

NestedValueProviders meet this need

```
public static void main(String[] args) {  
    pipeline  
        .apply(Create.of(1, 2, 3).withCoder(BigEndianIntegerCoder.of()));  
        // Write to the computed complete file path  
        .apply("Output+Nums" TextIO.write().to(NestedValueProvider.of(  
            options.getFileName(),  
            new SerializableFunction<String, String>() {  
                @Override  
                public String apply(String file) {  
                    return "gs://bucket/" + file;  
                }  
            }  
        ))));  
    pipeline.run();  
}
```



# Template Metadata

- Located at the same directory, named <template\_name>\_metadata

```
{
  "name": "WordCount",
  "description": "An example pipeline that counts words in the input file.",
  "parameters": [{
    "name": "inputFile",
    "label": "Input Cloud Storage File(s)",
    "help_text": "Path of the file pattern glob to read from.",
    "regexes": ["^gs://[/\\[^\\n\\r]+$"],
    "is_optional": true
  },
  {
    "name": "output",
    "label": "Output Cloud Storage File Prefix",
    "help_text": "Path and filename prefix for writing output files. ex: gs://MyBucket/counts",
    "regexes": ["^gs://[/\\[^\\n\\r]+$"]
  }
]
```

# Agenda

---

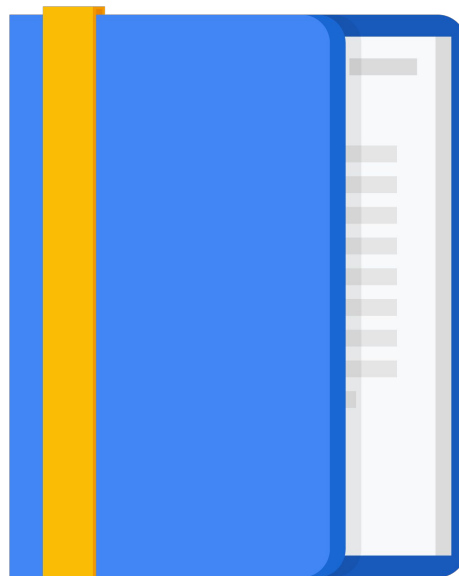
Cloud Dataflow

Why customers value Dataflow

Dataflow Pipelines

Dataflow Templates

Dataflow SQL



# Cloud Dataflow SQL lets you use SQL queries to develop and run Cloud Dataflow jobs from the BigQuery web UI

## Query editor

```
1  SELECT
2      sr.sales_region,
3      TUMBLE_START("INTERVAL 15 SECOND") AS period_start,
4      SUM(tr.payload.amount) as amount
5  FROM pubsub.topic.`dataflow-sql`.transactions AS tr
6      INNER JOIN bigquery.table.`dataflow-sql`.dataflow_sql_dataset.us_state_salesregions AS sr
7      ON tr.payload.state = sr.state_code
8  GROUP BY
9      sr.sales_region,
10     TUMBLE(tr.event_timestamp, "INTERVAL 15 SECOND")
```



Valid.

Cloud Dataflow engine alpha



Create Cloud Dataflow job



More ▾